

# AutoSA: A Polyhedral Compiler for High-Performance Systolic Arrays on FPGAs

**Jie Wang**, Licheng Guo, Jason Cong

University of California, Los Angeles

# Outline

- **A Brief Background**
- Details about AutoSA Compilation Flow
- Demos
  - Matrix Multiplication
  - Convolutional Neural Network
- Auto-Tuning in AutoSA
- Using AutoBridge to Boost the Design Frequency

# AutoSA: A Polyhedral Compiler for High-Performance Systolic Arrays on FPGAs

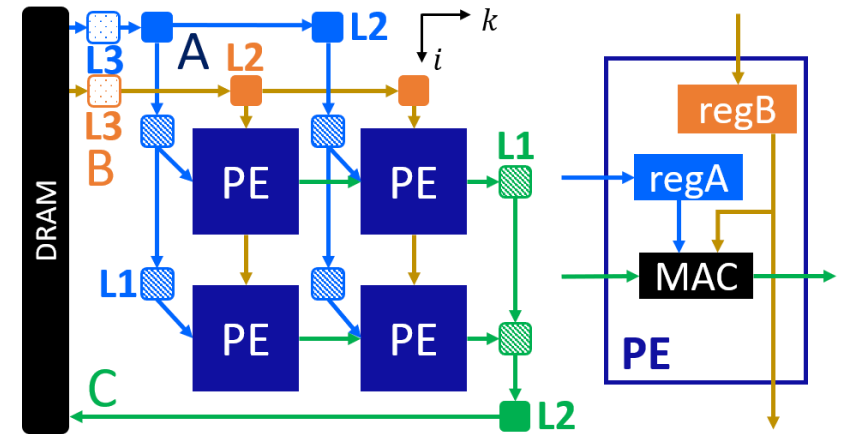
Wang, Jie, Licheng Guo, and Jason Cong. "AutoSA: A Polyhedral Compiler for High-Performance Systolic Arrays on FPGA." Proceedings of the 2021 ACM/SIGDA international symposium on Field-programmable gate arrays. 2021.

```
#pragma scop
for (int i = 0; i < M; ++i)
  for (int j = 0; j < N; ++j) {
S0:    C[i][j] = 0;
    for (int k = 0; k < K; ++k)
S1:      C[i][j] += A[i][k] * B[k][j];
  }
#pragma endscop
```

Input: C code



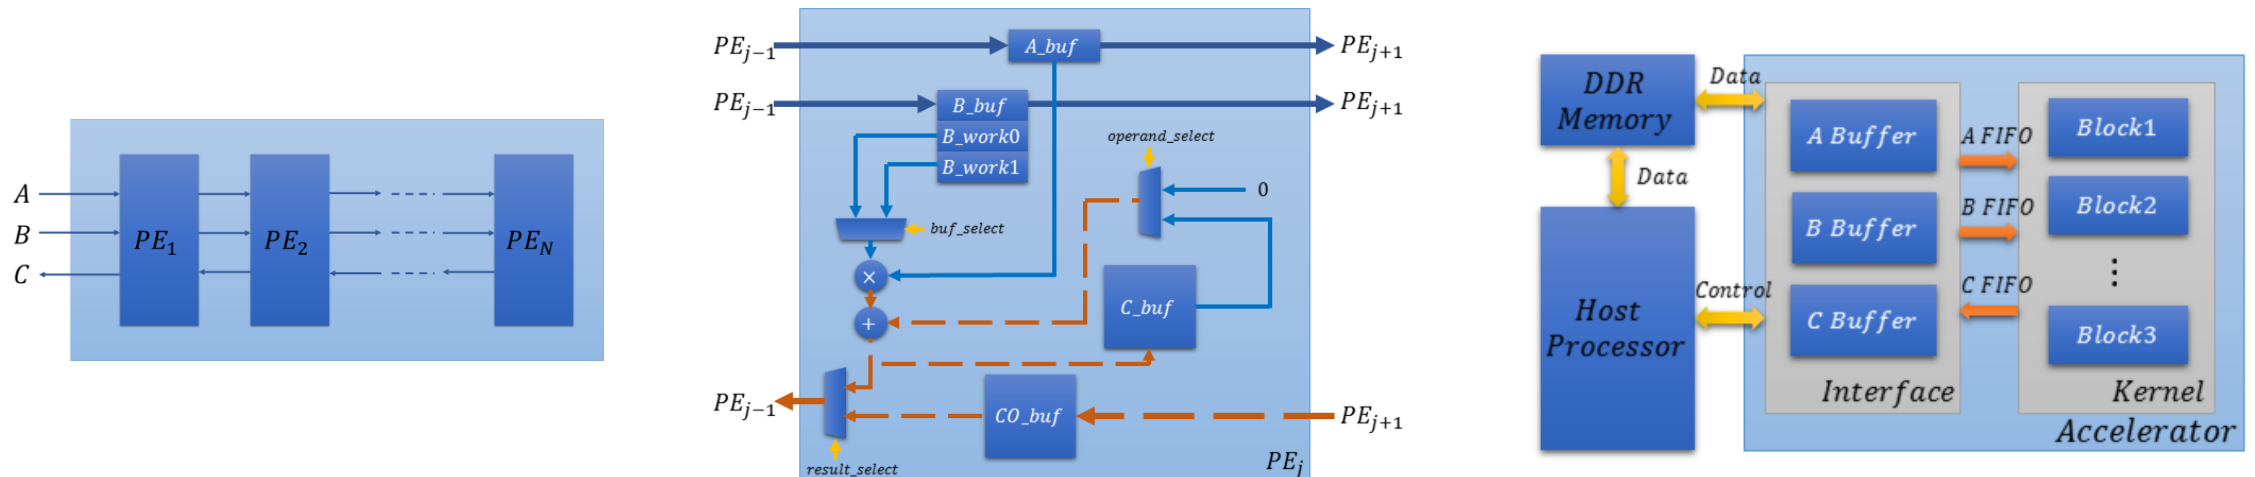
**AutoSA**



Output: Systolic array design in  
HLS C

# Starting Point

- 2014 Summer Internship at UCLA
- Time: ~5 months
- A manual-designed RTL-based 1D systolic array for MM
- 198.1 GFLOPs on Xilinx VC709



# Previous Work

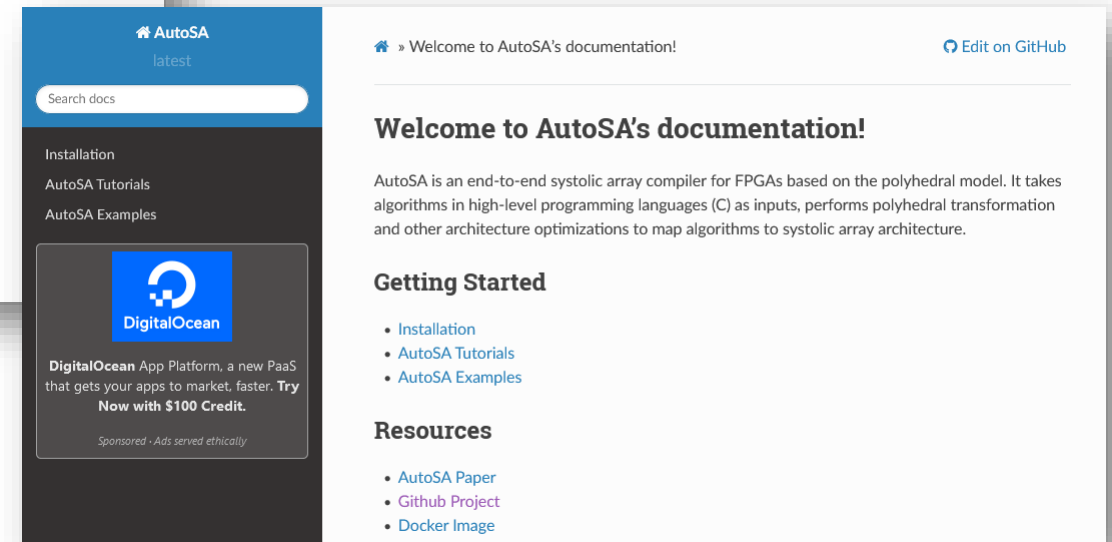
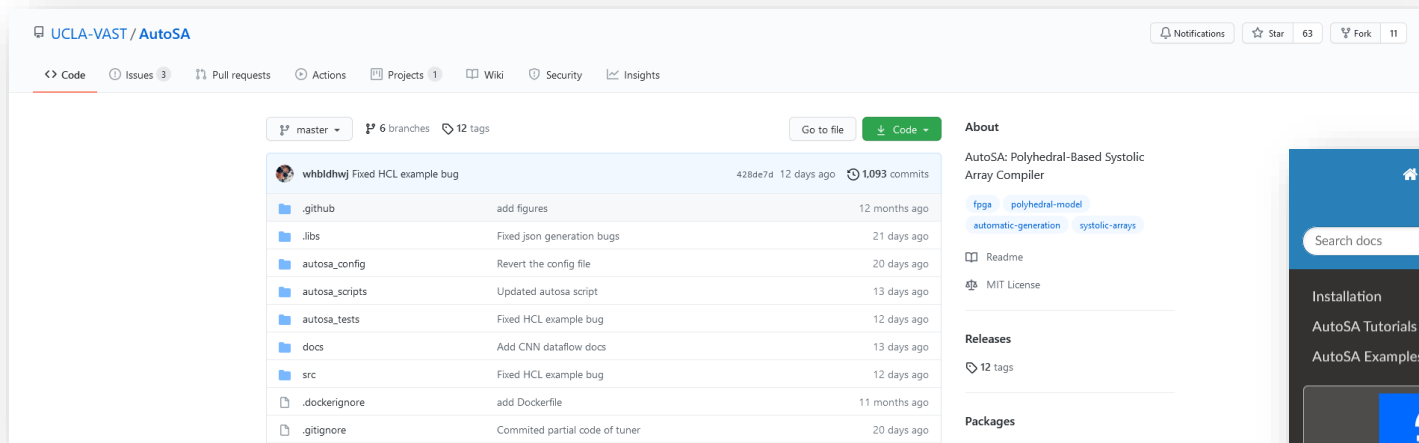
	MMAAlpha '01	Bondhugula et al. '07	PolySA '18	SuSy '20	AutoSA '21
<b>Generality</b>	Support applications as general as possible.				
Imperfectly Nested Loops	No	No	No	No	Yes
Multi-Statement	Yes	No	No	Yes	Yes
<b>Performance</b>	Generate designs with performance as high as possible.				
Latency Hiding	No	No	Auto	Semi-Auto	Auto
Double Buffering	No	No	Auto	Semi-Auto	Auto
<b>Productivity</b>	Shorten the development time as much as possible.				
Auto-Tuning	No	No	Yes	No	Yes

# Outline

- A Brief Background
- **Details about AutoSA Compilation Flow**
- Demos
  - Matrix Multiplication
  - Convolutional Neural Network
- Auto-Tuning in AutoSA
- Using AutoBridge to Boost the Design Frequency

# AutoSA is Open-Sourced

- Github: <https://github.com/UCLA-VAST/AutoSA>
- Document: <https://autosa.readthedocs.io/en/latest/>



# Current Capabilities

- Back-end
  - Xilinx HLS C (Mature, recommended)
  - Intel OpenCL (Experimental)
  - Catapult HLS C (Experimental)
- High-Performance designs on Xilinx platforms






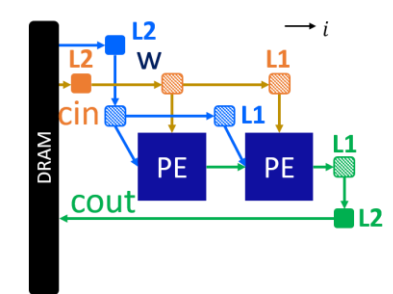





Matrix Multiplication	FP32	int16	Int8
MHz	300	250	300
TFLOPs	0.93	3.41	6.95

Board: Xilinx Alveo U250




# Current Capabilities

- **Architectural exploration** of different dataflows
- Example: 10 dataflows of CNN  
(<https://autosa.readthedocs.io/en/latest/examples/cnn.html>)



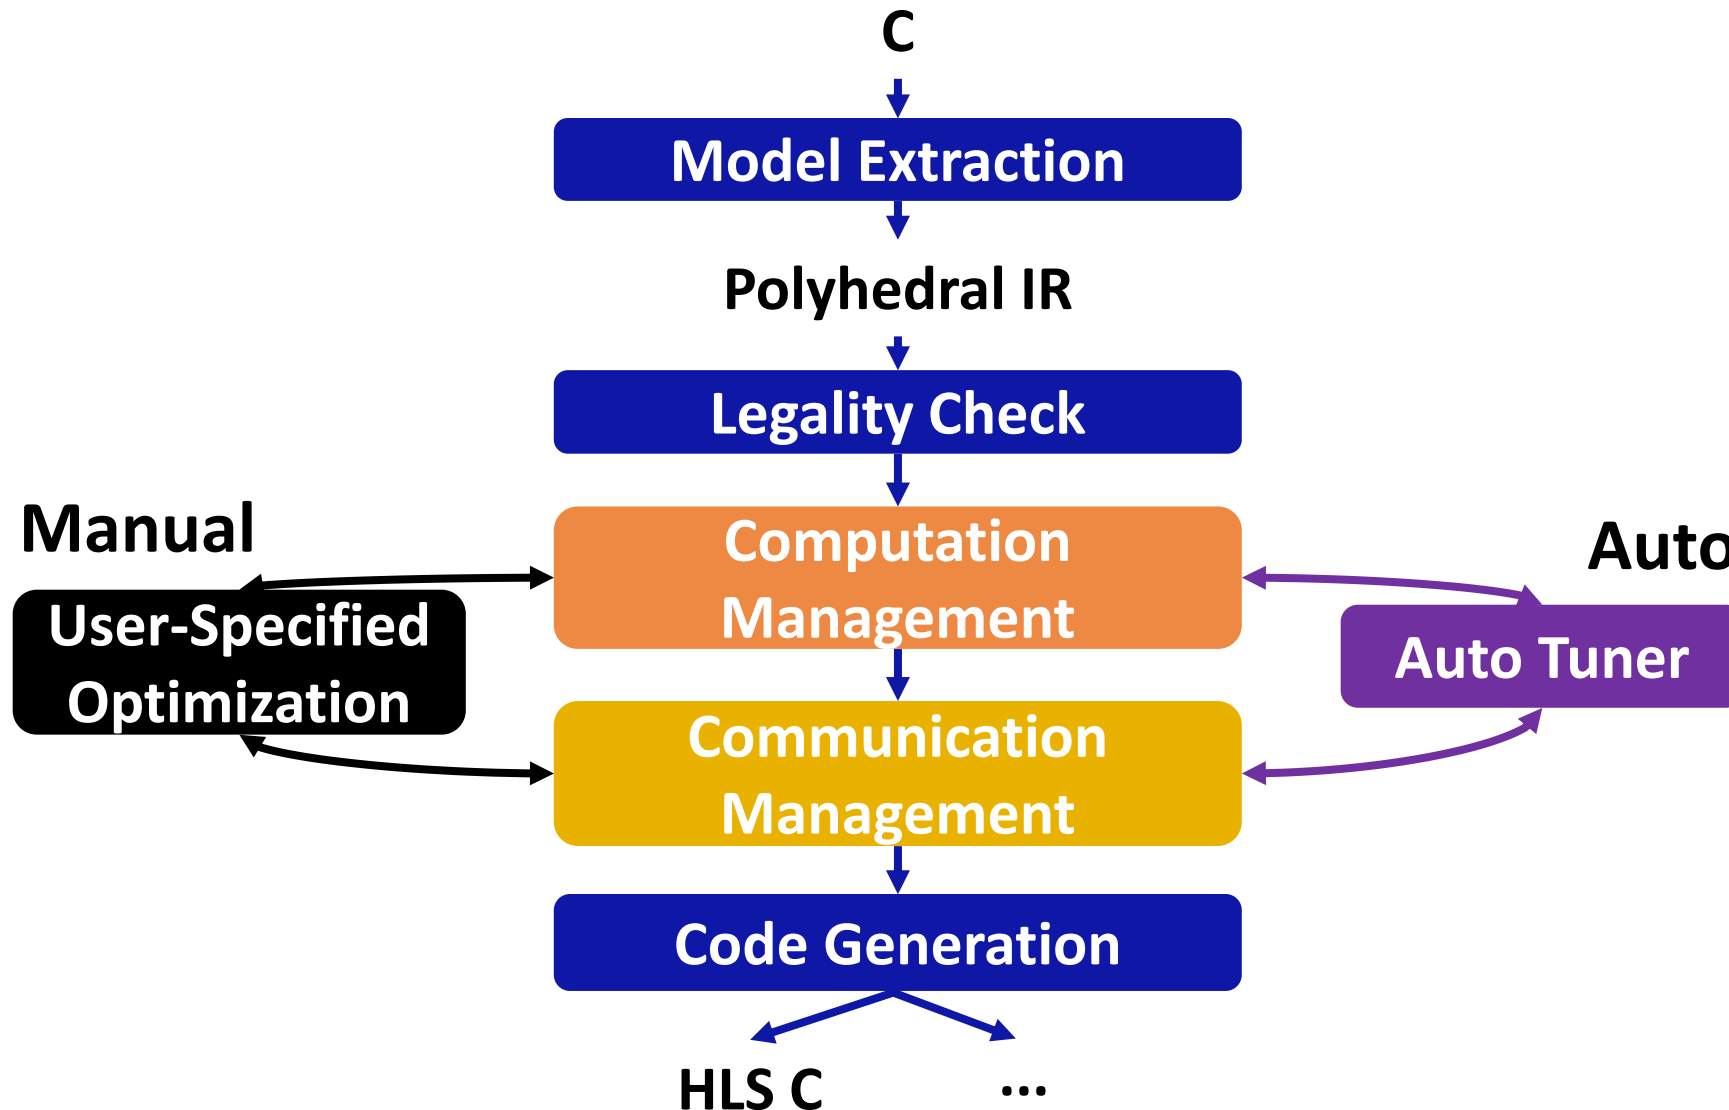
# Current Capabilities

- Auto-tuning support



- Exhaustive search with pruning
  - mins to hours
- Genetic search
  - < 1 minute

# Compilation Flow




# Compilation Flow

## c

```
#pragma scop
for (int i = 0; i < M; ++i)
    for (int j = 0; j < N; ++j) {
S0:        C[i][j] = 0;
            for (int k = 0; k < K; ++k)
S1:        C[i][j] += A[i][k] * B[k][j];
    }
#pragma endscop
```

The input program of AutoSA


# Compilation Flow



## Polyhedral Model

- A **mathematical** compilation framework for **loop nest optimization**.
- Applicable to a **subset** of general loop programs (static control of parts (SCoP)).
- **Complete** and **robust** toolchains.

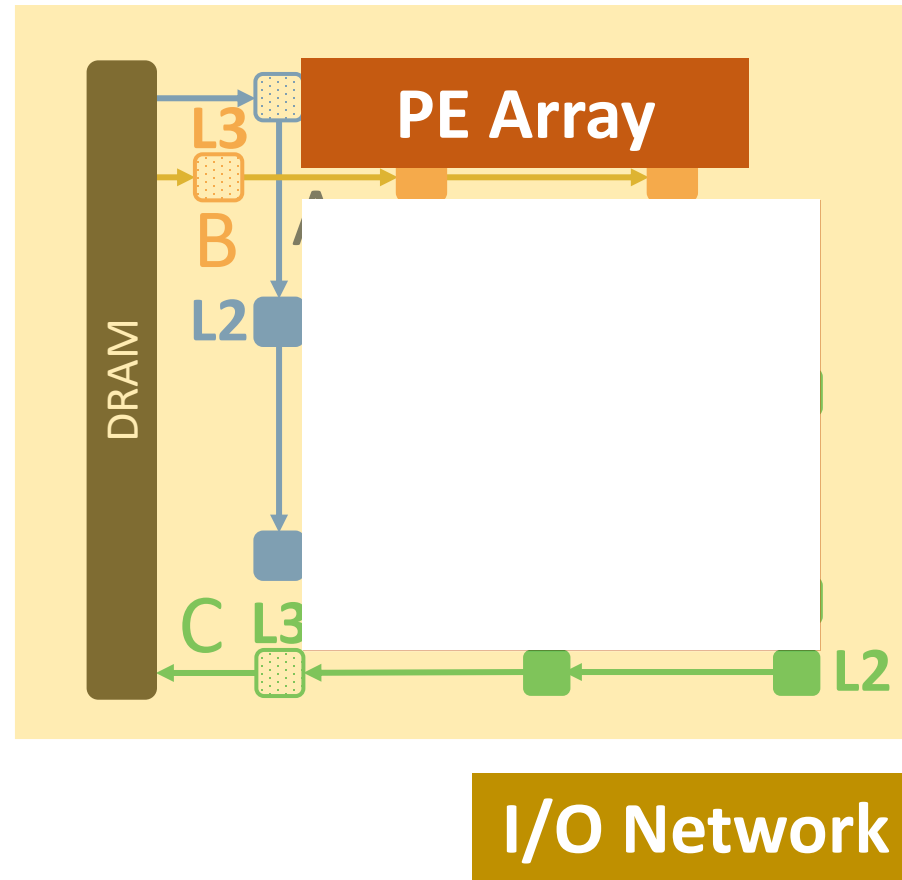
# Compilation Flow



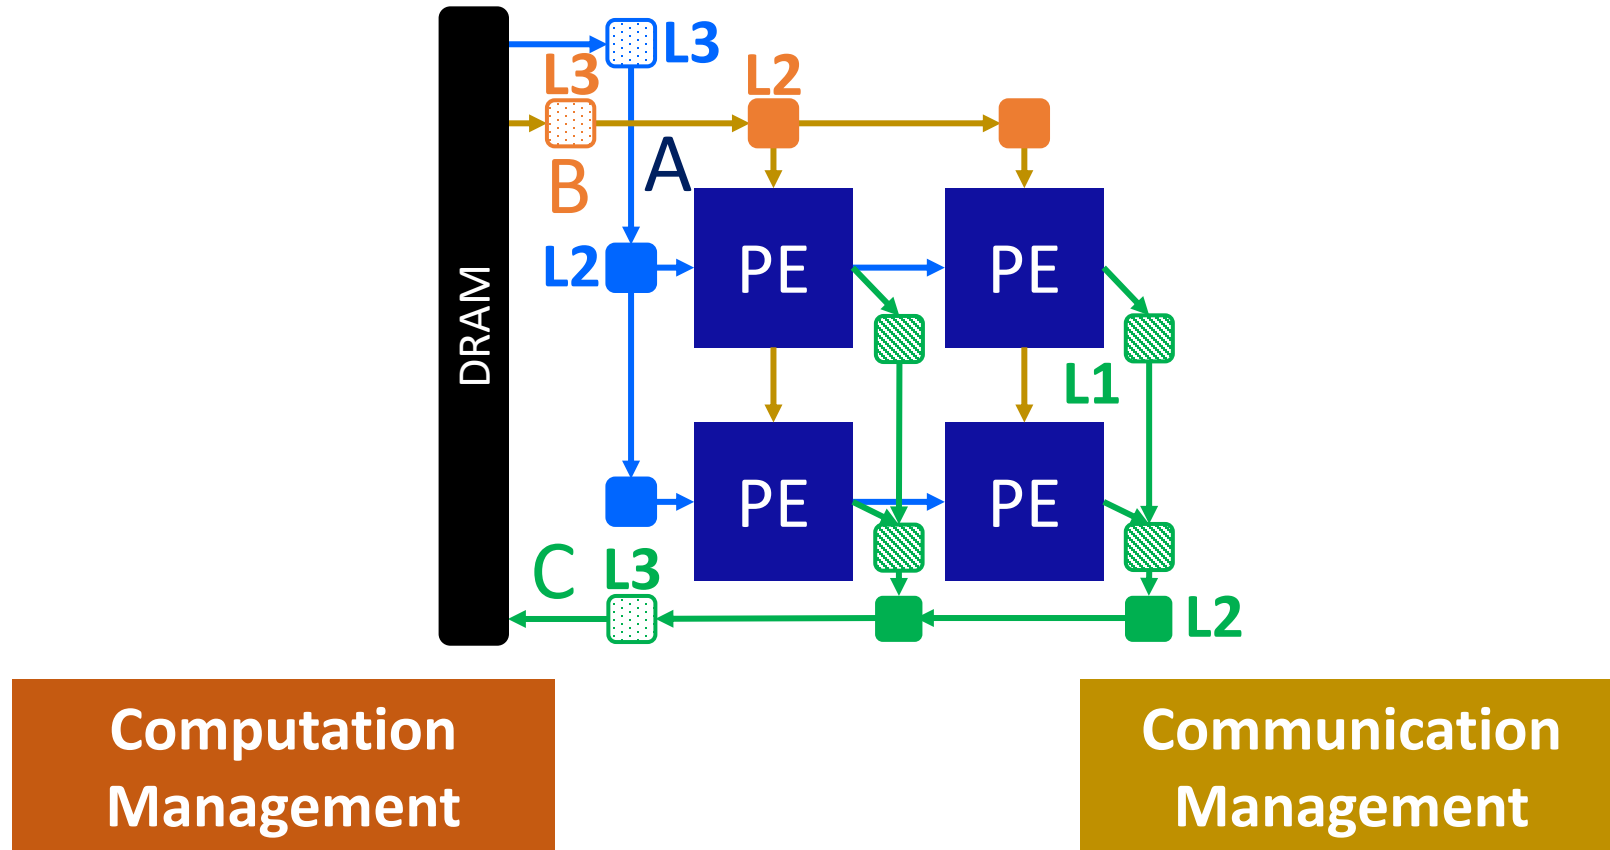
## Legality Check

- All dependences are **uniform** dependences (with constant dependence distance)  
➡ **Statically-scheduled** design
- Dependence distances on space loops (loops to be mapped to PE dimensions) to be **no greater than one**.  
➡ **Local** communication

# Compilation Flow

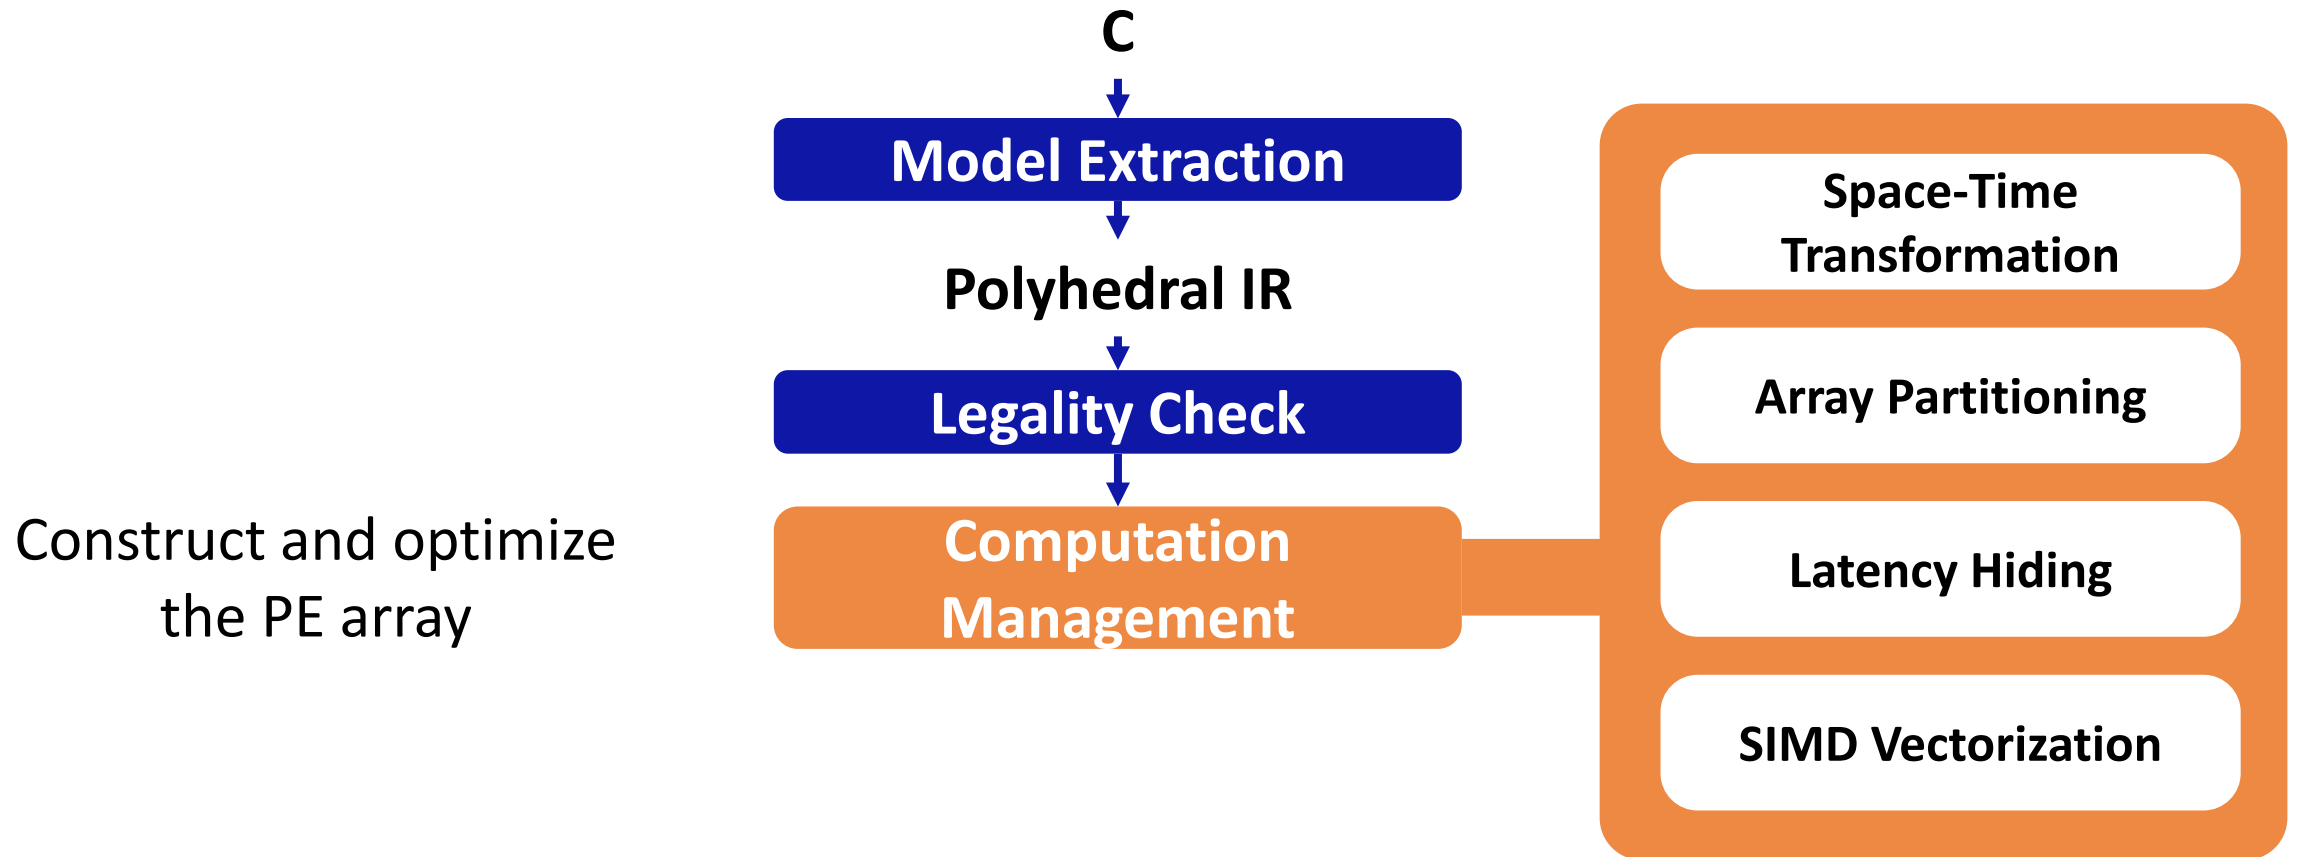


# Compilation Flow





# Compilation Flow



# Compilation Flow

C

**Model Extraction**

**Polyhedral IR**

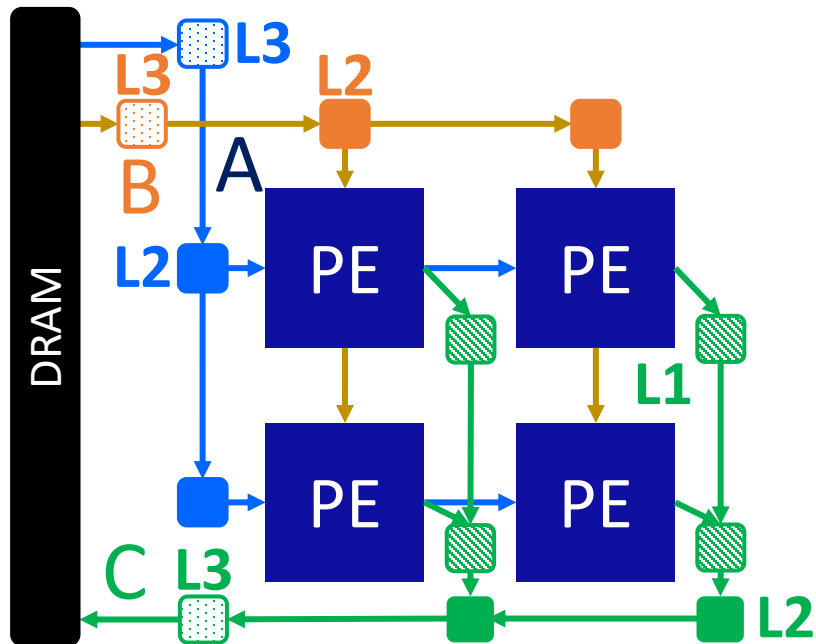
**Legality Check**

**Computation  
Management**

**Communication  
Management**

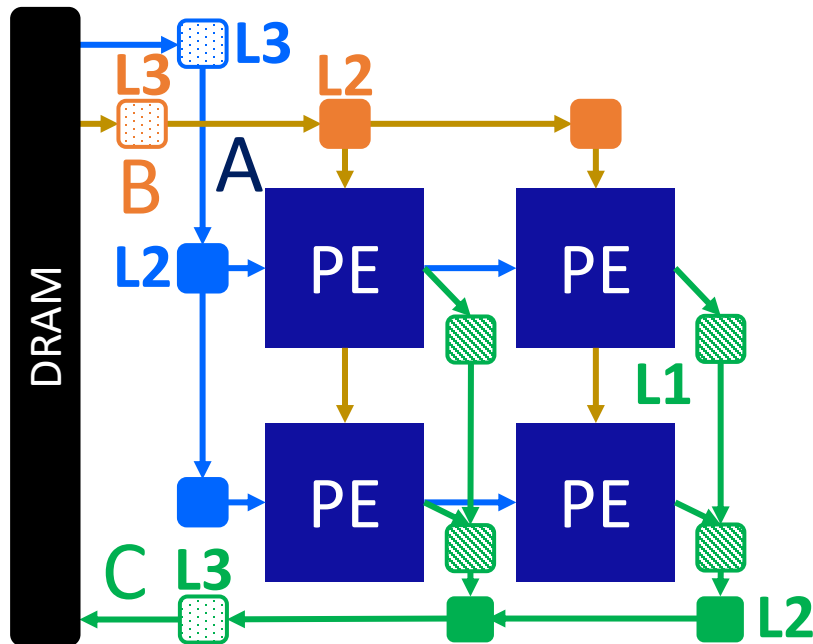
Construct and optimize  
the I/O network

# Communication Management



- **What** hardware modules to be generated for the I/O network?
- **How** to generate the I/O network?
- **How** to **optimize** the I/O network?

# Communication Management






**What** hardware modules to be generated for the I/O network?

- Inter-array communication modules
  - Data transfer logic inside PEs
  - PE interconnects
- Outer-array communication modules
  - I/O modules

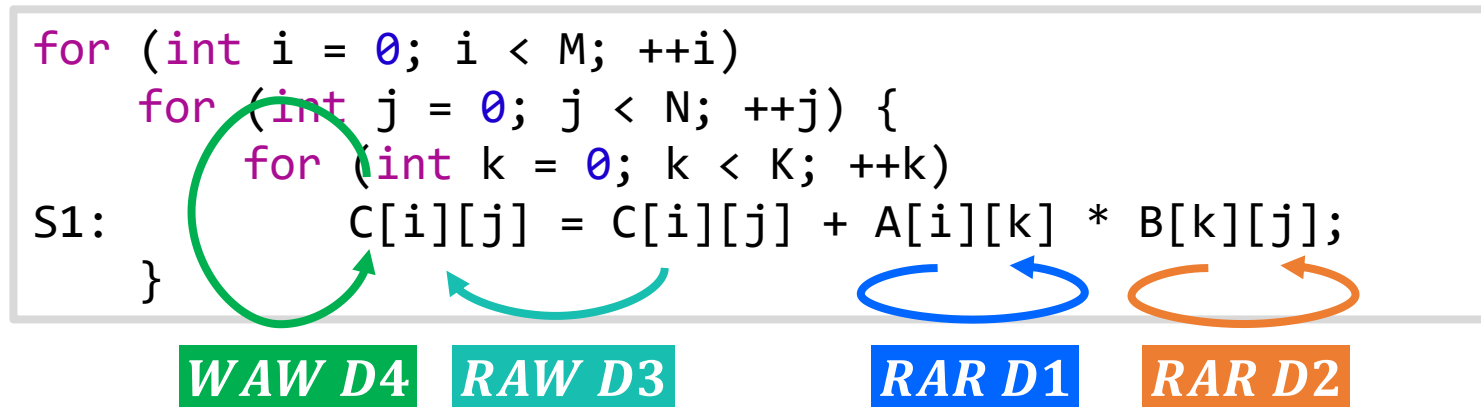
# Communication Management

**How** to generate the I/O network?

- **Key idea: Data communication** is determined by the data dependences.

Read dependences (Read-After-Read, RAR)		Read-only data
Flow dependences (Read-After-Write, RAW)		Intermediate data
Output dependences (Write-After-Write, WAW)		Computation results

# Communication Management



---

**Dependence**

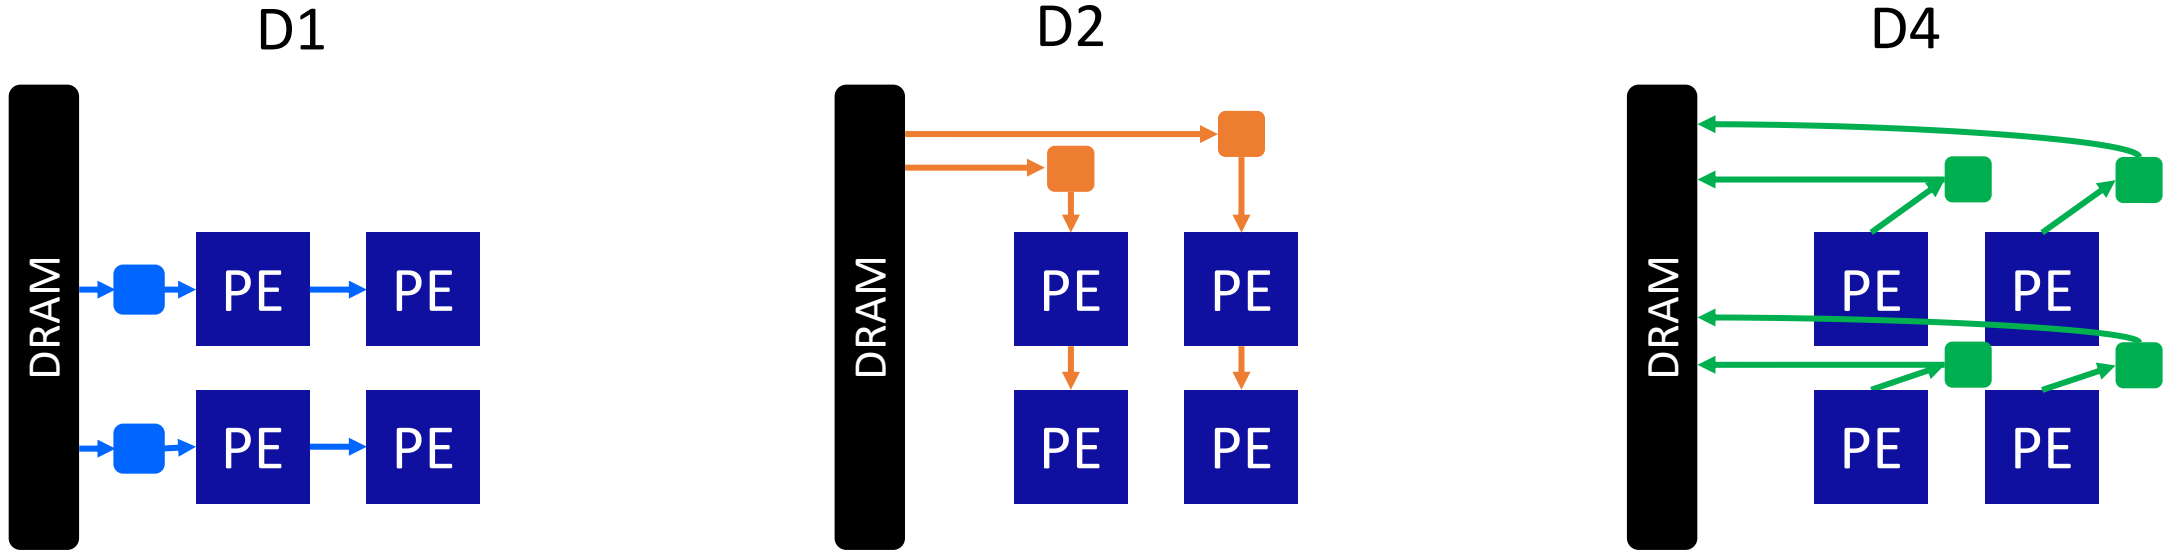
**Dependence Type**

**Array Access**

**Dependence  
Distance**

---

*We omit the statement of array initialization for the sake of brevity.*





Dependence	Dependence Type	Array Access	Dependence Distance
D1	Read (RAR)	$A[i][k]$	$(0,1,0)$
D2	Read (RAR)	$B[k][j]$	$(1,0,0)$
D3	Flow (RAW)	$C[i][j]$	$(0,0,1)$
D4	Output (WAW)	$C[i][j]$	$(0,0,1)$

*We omit the statement of array initialization for the sake of brevity.*

# Communication Management


## How to optimize the I/O network?

- I/O network localization
- Double buffering, data packing, etc.

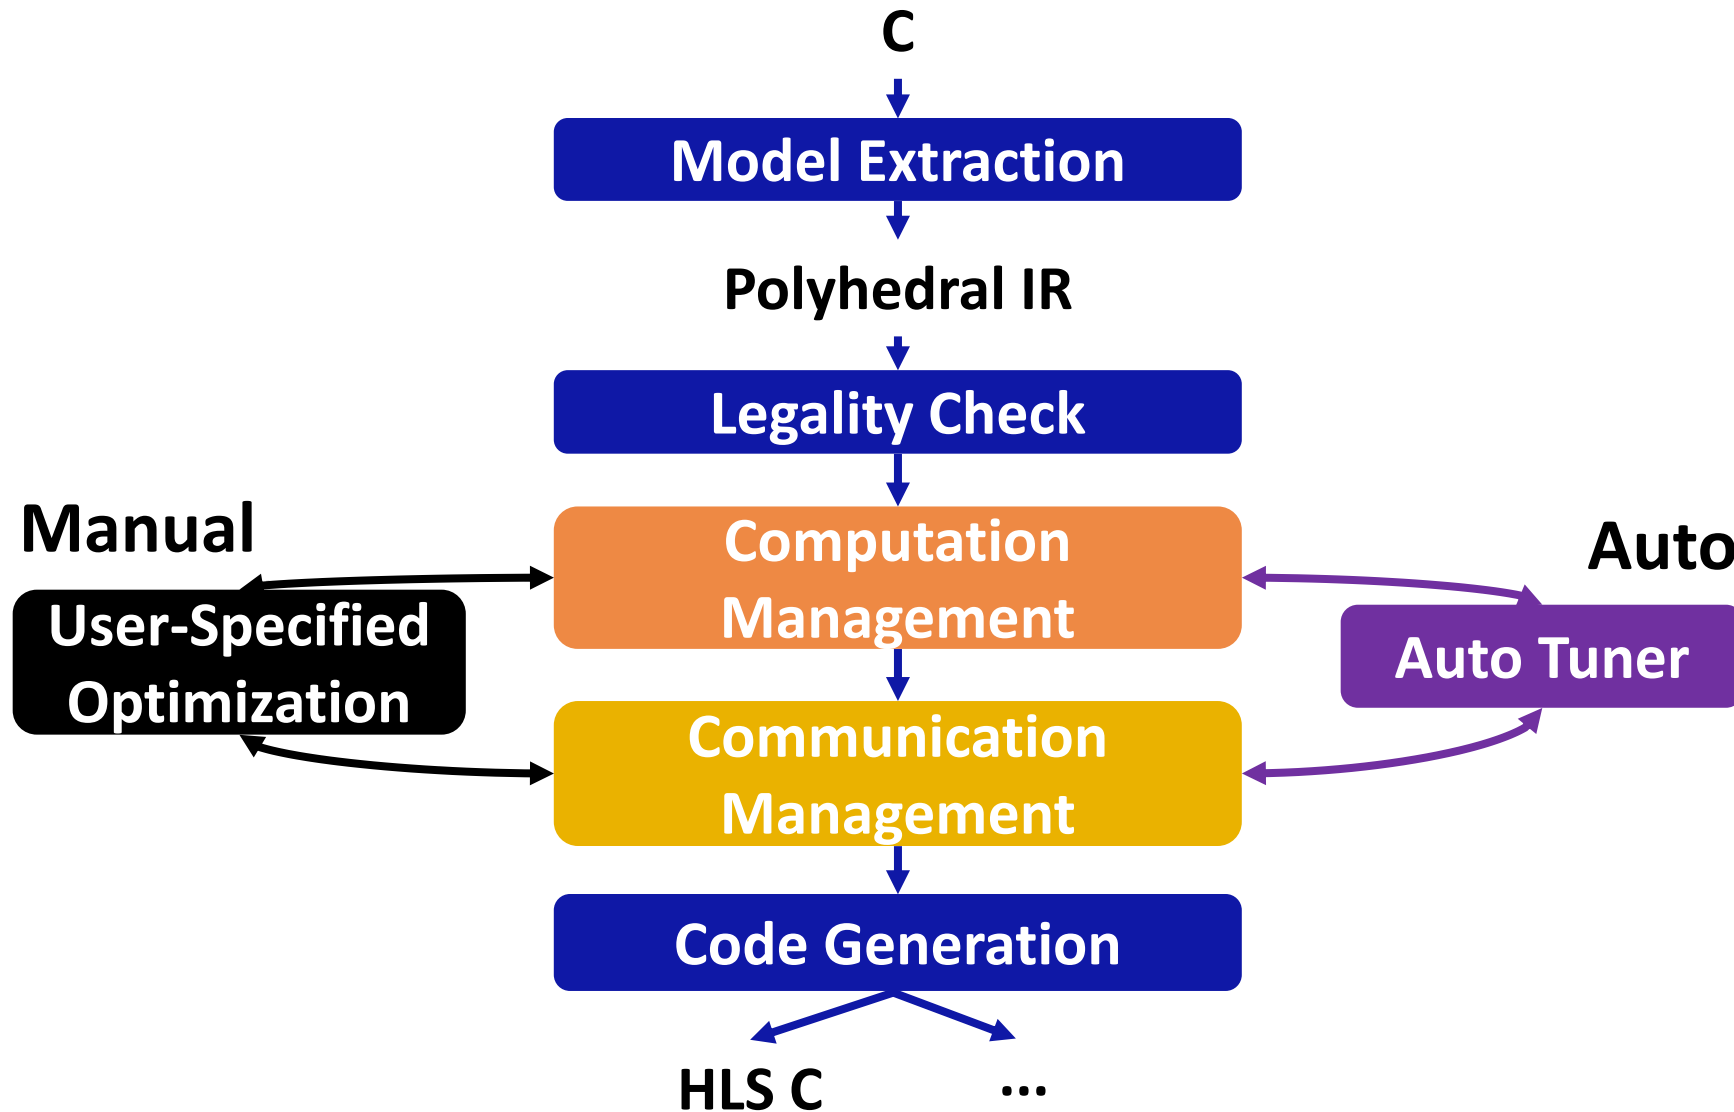




# Compilation Flow



# Compilation Flow



# Outline

- A Brief Background
- Details about AutoSA Compilation Flow
- **Demos**
  - Matrix Multiplication
  - Convolutional Neural Network
- Auto-Tuning in AutoSA
- Using AutoBridge to Boost the Design Frequency

# Demo 1: Matrix Multiplication

- Document:  
<https://autosa.readthedocs.io/en/latest/examples/mm.html>
- Example directory: `${AUTOSA_ROOT}/autosa_tests/mm`

# A Closer Look at Computation Management

- **Space-time mapping**: transforming the program to a systolic array with space-time mapping
- **Array partitioning**: partitioning the array into smaller sub-arrays to fit limited on-chip resource
- **Latency hiding**: permuting parallel loops inside to hide computation latency
- **SIMD vectorization**: vectorizing computation to amortize the PE control overheads

```
#pragma scop
for (int i = 0; i < M; ++i)
  for (int j = 0; j < N; ++j) {
S0:    C[i][j] = 0;
      for (int k = 0; k < K; ++k)
S1:    C[i][j] += A[i][k] * B[k][j];
  }
#pragma endscop
```



```
Space Loops
[ for (int i = 0; i < I; i++)
  [ for (int j = 0; j < J; j++) {
    [ for (int k = 0; k < K; k++) {
      C[i][j] += A[i][k] * B[k][j];
    }
  }
Time Loops
```

# A Closer Look at Computation Management

- **Space-time mapping:** transforming the program to a systolic array with space-time mapping
- **Array partitioning:** partitioning the array into smaller sub-arrays to fit limited on-chip resource
- **Latency hiding:** permuting parallel loops inside to hide computation latency
- **SIMD vectorization:** vectorizing computation to amortize the PE control overheads

**Space Loops**

```
for (int i = 0; i < I; i++)  
  for (int j = 0; j < J; j++) {  
    for (int k = 0; k < K; k++) {  
      C[i][j] += A[i][k] * B[k][j];  
    }  
  }
```


**Time Loops**



**Array Partitioning Loops**

```
for (int i = 0; i < M/4; i++)  
  for (int j = 0; j < N/4; j++)  
    for (int k = 0; k < K/4; k++)  
      for (int ii = 0; ii < 4; ii++)  
        for (int jj = 0; jj < 4; jj++)  
          for (int kk = 0; kk < 4; kk++)  
            S1;
```

**Sub-Array Loops**




# A Closer Look at Computation Management

- **Space-time mapping:** transforming the program to a systolic array with space-time mapping
- **Array partitioning:** partitioning the array into smaller sub-arrays to fit limited on-chip resource
- **Latency hiding:** permuting parallel loops inside to hide computation latency
- **SIMD vectorization:** vectorizing computation to amortize the PE control overheads

## Array Partitioning Loops

```
for (int i = 0; i < M/4; i++)  
  for (int j = 0; j < N/4; j++)  
    for (int k = 0; k < K/4; k++)  
      for (int ii = 0; ii < 4; ii++)  
        for (int jj = 0; jj < 4; jj++)  
          for (int kk = 0; kk < 4; kk++)  
            S1;
```

## Sub-Array Loops




4x4 array



## Latency Hiding Loops

```
for (int i = 0; i < M/4; i++)  
  for (int j = 0; j < N/4; j++)  
    for (int k = 0; k < K/4; k++)  
      for (int ii = 0; ii < 2; ii++)  
        for (int jj = 0; jj < 2; jj++)  
          for (int kk = 0; kk < 4; kk++)  
            for (int iii = 0; iii < 2; iii++)  
              #pragma HLS PIPELINE II=1  
                for (int jjj = 0; jjj < 2; jjj++)  
                  S1;
```



2x2 array

# A Closer Look at Computation Management

- **Space-time mapping**: transforming the program to a systolic array with space-time mapping
- **Array partitioning**: partitioning the array into smaller sub-arrays to fit limited on-chip resource
- **Latency hiding**: permuting parallel loops inside to hide computation latency
- **SIMD vectorization**: vectorizing computation to amortize the PE control overheads

```
for (int i = 0; i < M/4; i++)
  for (int j = 0; j < N/4; j++)
    for (int k = 0; k < K/4; k++)
      for (int ii = 0; ii < 2; ii++)
        for (int jj = 0; jj < 2; jj++)
          for (int kk = 0; kk < 4; kk++)
            for (int iii = 0; iii < 2; iii++)
              #pragma HLS PIPELINE II=1
              for (int jjj = 0; jjj < 2; jjj++)
                S1;
```

Latency Hiding Loops



```
for (int i = 0; i < M/4; i++)
  for (int j = 0; j < N/4; j++)
    for (int k = 0; k < K/4; k++)
      for (int ii = 0; ii < 2; ii++)
        for (int jj = 0; jj < 2; jj++)
          for (int kk = 0; kk < 2; kk++)
            for (int iii = 0; iii < 2; iii++)
              #pragma HLS PIPELINE II=1
              for (int jjj = 0; jjj < 2; jjj++)
                #pragma HLS UNROLL
                for (int kkk = 0; kkk < 2; kkk++)
                  S1;
```

SIMD Loops



# Demo 2: CNN

- Document:  
<https://autosa.readthedocs.io/en/latest/examples/cnn.html>
- Example directory: `${AUTOSA_ROOT}/autosa_tests/cnn`

# Auto-Tuning

Generality



Program with restrictions

- Rectangular iteration domain

Arbitrary program

**Genetic Search**

**Exhaustive Search  
with Pruning**

< 1 minute

Minutes to hours





Convergence Time

# Auto-Tuning

- Exhaustive search
  - [https://autosa.readthedocs.io/en/latest/tutorials/auto\\_tuning\\_exhaustive.html](https://autosa.readthedocs.io/en/latest/tutorials/auto_tuning_exhaustive.html)
- Genetic search
  - [https://autosa.readthedocs.io/en/latest/tutorials/auto\\_tuning\\_genetic.html](https://autosa.readthedocs.io/en/latest/tutorials/auto_tuning_genetic.html)
- Demo:
  - Use genetic search for matrix multiplication on Alveo U250

# Getting High-Frequency Systolic Array Designs with Floorplanning

Default place/router make suboptimal choices.  
Design Freq.:  
**216MHz**



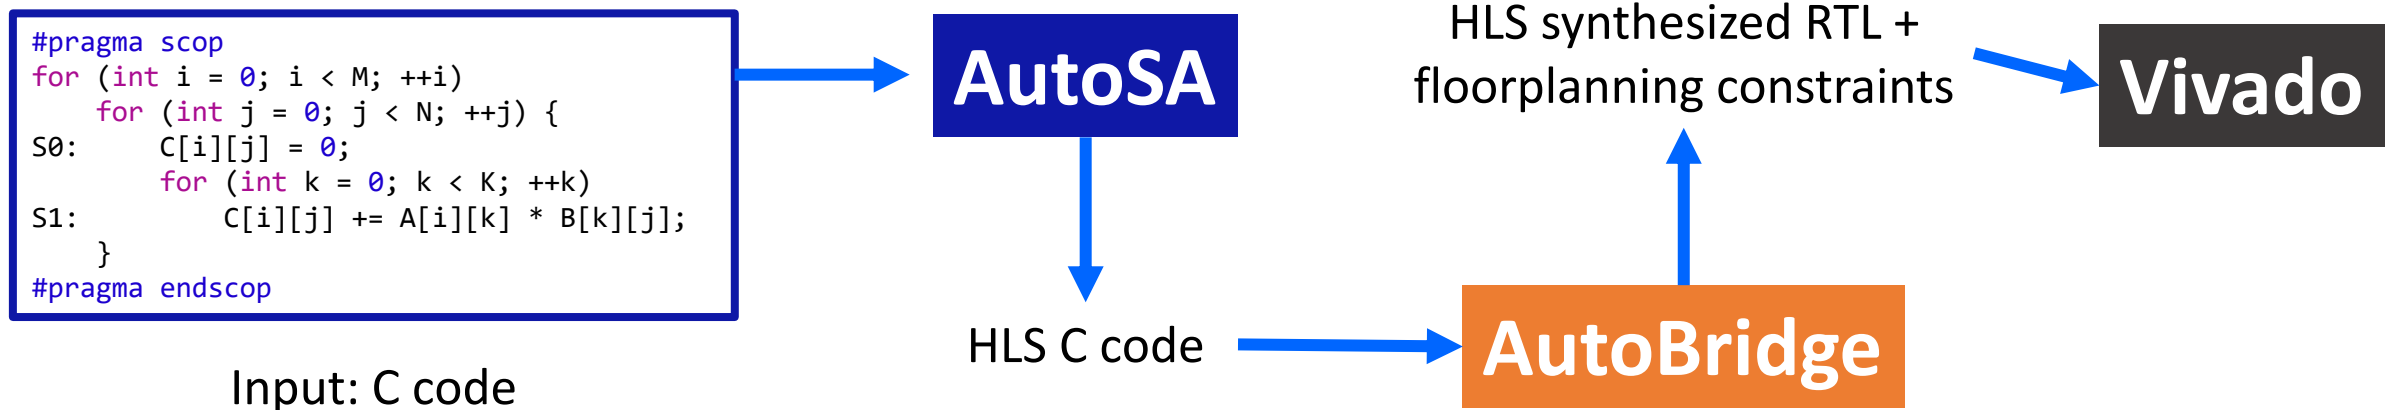
Floorplanning-guided approach spreads the logic across chip to reduce local congestion.  
Design Freq.:  
**329MHz**

Example: Improve the CNN systolic array design frequency using AutoBridge

# Automated Interconnect Pipelining for Multi-Die FPGA Designs

Guo, Licheng, et al. "**AutoBridge**: Coupling Coarse-Grained Floorplanning and Pipelining for High-Frequency HLS Design on Multi-Die FPGAs." (**FPGA 2021 Best Paper**)

- AutoBridge is an HLS-based automatic floorplanning flow.
- Improved the average frequency of 43 designs from 147MHz to 297MHz with a negligible area overhead.
- Github: <https://github.com/Licheng-Guo/AutoBridge>




# More Details

- An example of using AutoBridge with AutoSA
- [https://autossa.readthedocs.io/en/latest/tutorials/auto\\_bridge.html](https://autossa.readthedocs.io/en/latest/tutorials/auto_bridge.html)

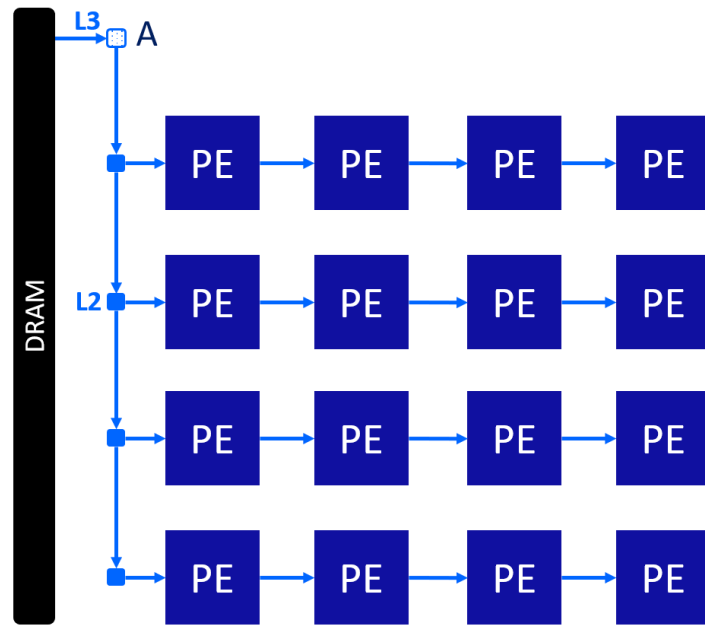
# Other Features

- Support for structural sparsity
  - [https://autosa.readthedocs.io/en/latest/tutorials/structural\\_sparsity.html](https://autosa.readthedocs.io/en/latest/tutorials/structural_sparsity.html)

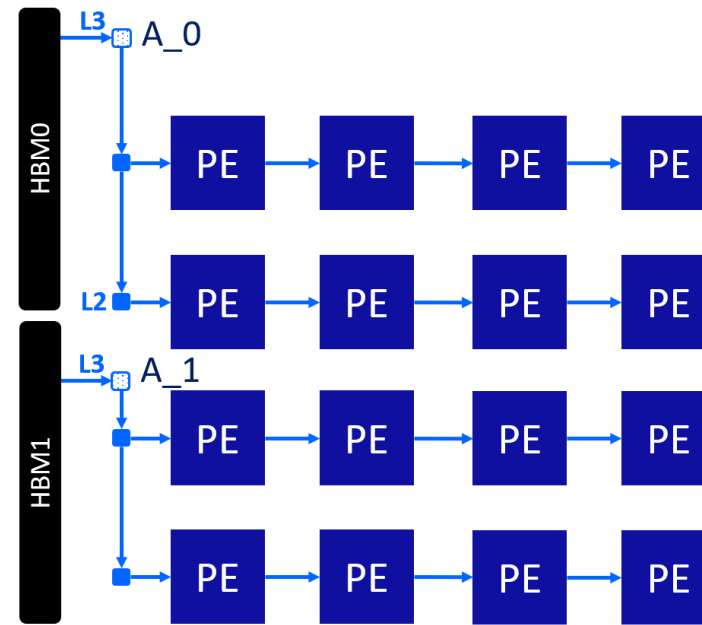


# Other Features

- Support for HBM
  - [https://autosa.readthedocs.io/en/latest/examples/mm\\_hbm.html](https://autosa.readthedocs.io/en/latest/examples/mm_hbm.html)



Array Using Single DDR/HBM Bank



Array Using Multiple DDR/HBM Banks



# How Can This Tool Benefit the Community?

- Understand the mechanisms of systolic array architecture, HLS design methodology, and polyhedral compilation
  - UCLA CS259
- Integrate AutoSA-generated designs into the accelerator framework
  - FlexCNN, HeteroCL
- Architecture studies based on an open-sourced and performant baseline
  - Architecture trade-offs for applications (dataflows, array configs, ...)
  - Programmable systolic array
  - Comparison against CGRA
  - ...

# Conclusion

- AutoSA, an end-to-end systolic array compiler for FPGA
- Goal: maximize generality, performance, and productivity when generating systolic arrays.
- Open-source and under active development. Welcome for any feedbacks and contribution!
- Any other questions?
- [jiewang@cs.ucla.edu](mailto:jiewang@cs.ucla.edu)

# User-Specified Optimizations

Option	Explanation	Example
array-contraction	Apply array contraction to reduce the memory usage.	True/False
axi-stream	Generate AXI stream interface.	True/False
block-sparse	Use block sparsity.	True/False
block-sparse-ratio	Block sparsity ratio.	[kernel[]->A[2,4]] Matrix A is sparse, every 4 elements are grouped into a vector, with 2 non-zero elements.
data-pack	Enable data-packing.	True/False
data-pack-sizes	Specifies the maximal data-packing factors at each I/O level.	[kernel[]->data_pack[8,32,64]]
double-buffer	Enable double buffering.	True/False
double-buffer-assign	Assign specific arrays to be double-buffered.	[kernel[]->A[]]

# User-Specified Optimizations

Option	Explanation	Example
explore-loop-permute	Explore loop permutation in the step of array partitioning.	True/False
fifo-depth	FIFO depth.	2
hbm	Use multi-port DRAM/HBM	True/False
hbm-port-num	Specify HBM port number per array.	2
hls	Generate Xilinx HLS host.	True/False
host-serialize	Serialize/deserialize the host data.	True/False
io-module-embedding	Embed I/O modules inside PEs if possible.	True/False
loop-infinite	Apply loop infinite optimization (Intel OpenCL only)	True/False

# User-Specified Optimizations

Option	Explanation	Example
local-reduce	Generate non-output-stationary array with local reduction.	True/False
reduce-op	Reduction operator.	“+”
max-sa-dim	Maximal systolic array dimension.	2
sa-sizes	Computation management parameters.	[kenrel[]->array_part[4,4]]